

Concepts in
Multicore
Programming

Lecture 7

Stencil Computations*

Charles E. Leiserson

February 25, 2010

*Slides courtesy of Matteo Frigo

Heat diffusion

1D heat diffusion equation:

$u(t, x)$: temperature at time t at position x .

$$\frac{\partial u}{\partial t} = \frac{\partial^2 u}{\partial x^2} .$$

Finite difference approximation:

$$\begin{aligned} \frac{\partial u}{\partial x}(t, x) &\approx \frac{u(t, x + \Delta x/2) - u(t, x - \Delta x/2)}{\Delta x} \\ \frac{\partial^2 u}{\partial x^2}(t, x) &\approx \frac{(\partial u / \partial x)(t, x + \Delta x/2) - (\partial u / \partial x)(t, x - \Delta x/2)}{\Delta x} \\ &\approx \frac{u(t, x + \Delta x) - 2u(t, x) + u(t, x - \Delta x)}{(\Delta x)^2} . \end{aligned}$$

3-point stencil

Finite differences for the heat diffusion equation:

$$\frac{u(t+1, x_i) - u(t, x_i)}{\Delta t} = \frac{u(t, x_{i-1}) - 2u(t, x_i) + u(t, x_{i+1}))}{(\Delta x)^2} .$$


Simple implementation:

```
for (t = 0; t < T; ++t) {           /* time loop */
    u[(t+1)%2][0] = left_boundary();
    for (i = 1; i < N - 1; ++i)      /* space loop */
        u[(t+1)%2][i] =
            kernel(u[t%2][i-1], u[t%2][i], u[t%2][i+1]);
    u[(t+1)%2][N - 1] = right_boundary();
}

double kernel(ui-1, ui, ui+1)
{
    return ui +  $\frac{\Delta t}{(\Delta x)^2}$  * (ui-1 - 2*ui + ui+1);
}
```

3-point stencil on a cache

```
for (t = 0; t < T; ++t) {           /* time loop */
    u[(t+1)%2][0] = left_boundary();
    for (i = 1; i < N - 1; ++i)      /* space loop */
        u[(t+1)%2][i] =
            kernel(u[t%2][i-1], u[t%2][i], u[t%2][i+1]);
    u[(t+1)%2][N - 1] = right_boundary();
}
```



If array u is larger than the cache, the number of misses is proportional to the number of accesses.


Cache oblivious algorithm for 3-point stencil

Recursively traverse trapezoidal regions of spacetime points (t, x) such that:

$$t_0 \leq t < t_1$$

$$x_0 + \dot{x}_0(t - t_0) \leq x < x_1 + \dot{x}_1(t - t_0)$$


$$\dot{x}_i \in \{-1, 0, 1\}$$



Base case


If height = 1, compute all spacetime points in the trapezoid.

Any order of computation is valid, because these points do not depend upon each other.



Space cut


If width $\geq 2 \cdot$ height, cut the trapezoid with a line of slope -1 through the center.



Traverse first the trapezoid on the left, then the one on the right.

Time cut

If $\text{width} < 2 \cdot \text{height}$, cut the trapezoid with a horizontal line through the center.



Traverse the bottom trapezoid first, then the top one.


C implementation

```
void trapezoid(int t0, int t1, int x0, int  $\dot{x}_0$ , int x1, int  $\dot{x}_1$ )
{
    int  $\Delta t$  = t1 - t0;
    if ( $\Delta t$  == 1) {
        int x;
        for (x = x0; x < x1; ++x)
            kernel(t0, x);
    } else if ( $\Delta t$  > 1) {
        if (2 * (x1 - x0) + ( $\dot{x}_1$  -  $\dot{x}_0$ ) *  $\Delta t$  >= 4 *  $\Delta t$ ) {
            int xm = (2 * (x0 + x1) + (2 +  $\dot{x}_0$  +  $\dot{x}_1$ ) *  $\Delta t$ ) / 4;
            trapezoid(t0, t1, x0,  $\dot{x}_0$ , xm, -1);
            trapezoid(t0, t1, xm, -1, x1,  $\dot{x}_1$ );
        } else {
            int s =  $\Delta t$  / 2;
            trapezoid(t0, t0 + s, x0,  $\dot{x}_0$ , x1,  $\dot{x}_1$ );
            trapezoid(t0 + s, t1, x0 +  $\dot{x}_0$  * s,  $\dot{x}_0$ , x1 +  $\dot{x}_1$  * s,  $\dot{x}_1$ );
        }
    }
}
```

Cache complexity of the stencil algorithm

When $\text{width} + \text{height} = \Theta(Z)$:

- ▶ number of cache misses = $O(\text{width} + \text{height})$.
- ▶ number of points = $\Theta(\text{width} \cdot \text{height})$.
- ▶ Algorithm guarantees that $\text{height} = \Theta(\text{width})$.
- ▶ Thus, $\text{height} = \Theta(Z)$, $\text{width} = \Theta(Z)$.
- ▶ Thus, number of cache misses = $\Theta(\text{number of points}/Z)$.



Demo

Simulation:

- ▶ $\Delta x = 95$.
- ▶ $\Delta t = 87$.
- ▶ $\dot{x}_0 = \dot{x}_1 = 0$.
- ▶ LRU cache.
- ▶ Line size = 4 points.
- ▶ Cache size = 4, 8, 16, or 32 cache lines.
- ▶ Cache miss latency = 10 cycles.