

Logic Programming — Impurities in Prolog

Steven Zeil

Apr. 1, 2001

Contents

1	Impurities in Prolog	2
1.1	Numerical Calculations	2
1.1.1	is versus =	2
1.1.2	is is not a pure relation.	2
1.1.3	Numeric Tests	3
1.2	Order Dependencies	3
1.2.1	Order and Grounding	3
1.2.2	Goal Ordering	3
1.2.3	Rule Ordering	5
1.3	Occurs Checks	6
1.4	Failure as Negation	6
1.4.1	Closed World Assumption	6
1.4.2	“not” in a Closed World	6
1.5	Cuts	7
1.5.1	Example of Cut	7
1.5.2	Applications of Cut	9
2	Programming in Prolog	10
2.1	Algorithmic Styles	10
2.1.1	Backtracking	10
2.1.2	Guess and Verify	10
2.2	Dare to Fail!	11
2.3	Open Lists	11
2.3.1	A Queue in Prolog	11
3	Closing thoughts	11

Logic Programming

1. Overview
2. Relations
3. Prolog Basics
4. Implementing LP
5. Impurities in Prolog
6. Programming in Prolog

1 Impurities in Prolog

1. Numerical Calculations
2. Order Dependencies
3. Occurs Checks
4. Failure as Negation
5. Cuts

1.1 Numerical Calculations

2 Evaluation of numerical expressions in Prolog is accomplished via `is`:

$\langle \text{simple-term} \rangle \text{ is } \langle \text{term} \rangle$

(A simple term is an atom, variable, or number).

For example,

```
A is 1+2*3.
X=1, Y=2, Z is X/Y.
3 is 1+2.
X=1, Y=2, 3 is X/Y.
```

The last example fails, but is still legal as a test.

1.1.1 is versus =

- `is` evaluates an expression
- `=` unifies expressions

Compare the results of

```
A is 1+2*3.
A = 1+2*3.
X=1, Y=2, Z is X/Y.
X=1, Y=2, Z = X/Y.
3 is 1+2.
3 = 1+2.
```

1.1.2 is is not a pure relation.

The term on the right hand side of an `is` must be **grounded**, i.e., all its variables must be bound to atoms, numbers, or to other grounded expressions.

- This is valid:
`X=1, Y=2, 0.5 is X/Y.`
- This is not:
`X=1, 0.5 is X/Y.`

It doesn't merely fail, it is an error that halts the execution.

1.1.3 Numeric Tests

- Prolog provides the usual range of relational operators:
 - $>$, $<$, $>=$, $=<$
 - Numerical equality test: $==$
 - Numerical not-equals test: $==$
- These do an implicit evaluation of their operands.
 - Therefore, like `is`, they operate on grounded terms only.

```

A is 1+2, A == 3.    /* yes */
1+2 == 3.           /* no */
A is 1+2, A > 0.     /* yes */
1+2 == B.           /* error */
A is 1+2, A < A+1.   /* yes */

```

1.2 Order Dependencies

The order in which we write goals and rules can affect

- speed
- meaning
- validity

of a Prolog program.

-
- Order and Grounding
 - Goal Ordering
 - Rule Ordering
-

1.2.1 Order and Grounding

An obvious example of order dependency arises with predicates that require grounded operands.

Compare, for example:

```

A=1+2, B=A.
B=A, A=1+2.
A=1+2, B==A.
B==A, A=1+2. /* error */

```

1.2.2 Goal Ordering

Given

```

prefix(A,B) :- append(A, _, B).
suffix(A,B) :- append(_, A, B).

```

use `';` to examine the possible solutions of:

```

prefix(X, [a,b,c]).
suffix([a], X).

```

Now try the combined queries:

```
prefix(X, [a,b,c]), suffix([a], X).
```

and

```
suffix([a], X), prefix(X, [a,b,c]).
```

The second query goes into an “infinite” recursion after `';`.

```
prefix(X, [a,b,c]), suffix([a], X).
```

These are placed on a list of pending goals.

Remaining goals: `prefix(X, [a,b,c]), suffix([a], X)`

The first goal is then visited.

The current goal unifies with the rule


```
prefix(A,B) :- append(A, _, B).
```

We then call (visit) the first of the remaining goals.


This first goal unifies with the rule

```
append([], B, B).
```

which unifies $\{X \rightarrow []\}$.




Again, we visit the first remaining goal.




This unifies with the rule

$\text{suffix}(A, B) :- \text{append}(_, A, B).$




Visit first remaining goal.



This unifies with the rule

$\text{append}([], B, B).$




This does not unify with either rule for append:


$\text{append}([], B, B).$

$\text{append}([X|A], B, [X|C]) :- \text{append}(A, B, C).$

So this goal fails and we redo the prior goal.




But suffix only has one rule, and we've already tried that, so this goal fails as well.




We move on to the 2nd rule for append, which unifies.

$\text{append}([Y|A], B, [Y|C]) :- \text{append}(A, B, C).$

with $\sigma = \{X \rightarrow [a|A], C \rightarrow [b, c]\}$




We now visit the first remaining goal.




The current goal unifies with the rule:

$\text{append}([], D, D).$



Visit the 1st remaining goal.



This unifies with

$\text{suffix}(A, B) :- \text{append}(_, A, B).$

And now for something completely different

Reverse the order of the goals:

`suffix([a], X), prefix(X, [a,b,c]).`

For the sake of brevity, we will omit the append subgoals.

Remaining goals: `suffix([a], X), prefix(X, [a,b,c])`

The first goal is then visited.

Remaining goals: `suffix([a], X), prefix(X, [a,b,c])`

The current goal eventually succeeds with $\sigma = \{X \rightarrow [a]\}$

Remaining goals: `prefix([a], [a,b,c])`

The current goal eventually succeeds and the system announces

$X = [a]$

Suppose we hit ';' to force a redo of the last goal. It then fails.

Remaining goals: `suffix([a], X), prefix(X, [a,b,c])`

On redo, the current goal eventually succeeds with $\sigma = \{X \rightarrow [A,a]\}$

Remaining goals: `prefix([A,a], [a,b,c])`

The prefix goal fails.

Remaining goals: `suffix([a], X), prefix(X, [a,b,c])`

On redo, the current goal eventually succeeds with $\sigma = \{X \rightarrow [B,C,a]\}$

Remaining goals: `prefix([B,C,a], [a,b,c])`

The prefix goal fails.

Remaining goals: `prefix(X, [a,b,c])`

On redo, the current goal eventually succeeds with $\sigma = \{X \rightarrow [D,E,F,a]\}$

Remaining goals: `suffix([D,E,F,a], [a,b,c])`

The prefix goal fails.

Remaining goals: `suffix([a], X), prefix(X, [a,b,c])`

On redo, the current goal eventually succeeds with $\sigma = \{X \rightarrow [G,H,I,J,a]\}$

This pattern continues forever, because

- there are an infinite number of lists that will satisfy `suffix([a], X)`,
- but none (except [a] will satisfy `prefix(X, [a,b,c])`.

1.2.3 Rule Ordering

- The Prolog engine attempts each unifying rule in the order they occur in the database.
- Changing that order alters the order of solutions attempted
- May affect meaning, speed, or termination of a query.

For example, we previously defined

`member(X, [X|_]).`

`member(X, [_|L]) :- member(X,L).`

The query `member(2,L)` would return successive solutions:

$L = [2|A]$

$L = [2,A|B]$

$L = [2,A,B|C]$


(Why isn't the first solution $L = [2]$? Did we miss a solution?)

If we reverse the order of those two rules


```
member2(X, [_|L]) :- member2(X,L).
member2(X, [X|_]).
```

The query `member(2,L)` goes into an infinite recursion.


```
member2(X, [_|L]) :- member2(X,L).
member2(X, [X|_]).
```




```
member2(X, [_|L]) :- member2(X,L).
member2(X, [X|_]).
```



```
member2(X, [_|L]) :- member2(X,L).
member2(X, [X|_]).
```



```
member2(X, [_|L]) :- member2(X,L).
member2(X, [X|_]).
```



and so on...

- One way to combat rule order dependency is to always list “base cases” before recursive cases.
 - Not fool-proof, though.
- Trying to anticipate these dependencies is hard because we need to consider all possible combinations of grounded/ungrounded operands.

1.3 Occurs Checks

- Prolog is essentially tree-oriented.
- If a variable `X` is unified with a term that contains `X`, this tree structure is violated.
 - `X=f(X)`

– `append(A, E, [a,b|E])`

- But for efficiency reasons, Prolog does not check for such cases.

1.4 Failure as Negation

Given the facts

```
student(john).
student(mary).
student(vladimir).
```

What is the response to the query: `student(brunhilda).?`

Why does it say “no”?

- Not because we *know* that Brunhilda is not a student,
- But because we cannot *prove* that she is

1.4.1 Closed World Assumption

Prolog operates under a **closed world assumption**:

- There are not facts or rules missing from our database’s description of the “world”.

1.4.2 “not” in a Closed World

- In a closed world, we should be able to prove anything that is true.
- So a natural definition of “not” would be:
 - `not(X)` succeeds if `X` fails
 - `not(X)` fails if `X` succeeds
- But this definition is very prone to order dependencies
- and can cause other surprises as well.

```
X=a, Y=b, X=Y.           /* no */
X=a, Y=b, not(X=Y).      /* yes */
X=a, not(X=Y), y=b.      /* yes */
X=a, not(X=Y).           /* no */
X=a, X=Y.                /* yes */
X=a, not(not(X=Y)).      /* yes */
```

Of course, if the closed world assumption does not hold, or other impurities prevent X from being proven, then $\text{not}(X)$ may mean something entirely different.

1.5 Cuts

Given rules

$$\begin{aligned} A &:- B_1, B_2, \dots \\ A &:- C_1, C_2, \dots \\ &\vdots \\ A &:- D_1, D_2, \dots, D_k, !, D_{k+1}, \dots \\ A &:- E_1, E_2, \dots \\ &\vdots \end{aligned}$$

The “!” is a **cut**.

When a cut is “solved” as a goal, we “commit” to the use of that rule.

- If a subsequent D_i term in that rule fails, then
 - this rule fails
 - no other rules for A will be attempted
 - * so the “parent” A goal will fail as well
- If the remaining terms in this rule succeed, but a later failure causes us to try to redo the parent A goal, then that parent goal fails.

1.5.1 Example of Cut


Given the rules:


$$\begin{aligned} a &:- b, f. \\ b &:- c, d. \\ b &:- d, !, e. \\ b &:- c. \\ c. \\ d. \\ e. \end{aligned}$$

what happens to the query: a ?


$$\begin{aligned} a &:- b, f. \\ b &:- c, d. \\ b &:- d, !, e. \end{aligned}$$


$$\begin{aligned} b &:- c. \\ c. \\ d. \\ e. \end{aligned}$$


$$\begin{aligned} a &:- b, f. \\ b &:- c, d. \\ b &:- d, !, e. \\ b &:- c. \\ c. \\ d. \\ e. \end{aligned}$$


$$\begin{aligned} a &:- b, f. \\ b &:- c, d. \\ b &:- d, !, e. \\ b &:- c. \\ c. \\ d. \\ e. \end{aligned}$$



The notation b_1 indicates that we have tried (are trying) rule 1.

$$\begin{aligned} a &:- b, f. \\ b &:- c, d. \\ b &:- d, !, e. \\ b &:- c. \\ c. \\ d. \\ e. \end{aligned}$$



$$\begin{aligned} a &:- b, f. \\ b &:- c, d. \\ b &:- d, !, e. \\ b &:- c. \\ c. \\ d. \\ e. \end{aligned}$$


- b has “succeeded” using rule 1
- but f fails, forcing a redo.


a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.




a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.



a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.




a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.




a :- b, f.
b :- c, d.
b :- d, !, e.

b :- c.
c.
d.
e.




a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.




The cut is “solved”. It forces us to commit to the current rule.

a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.



- The “remaining goals” list still contains the remaining portion of the 2nd b rule.
- But the already successful portion of that rule disappears from the backtrack list
- as does the b marker that was the head of the current rule.
- These are removed so that a later failure will not induce a redo of those terms.


a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.




```

a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.


```



```

a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.


```



```

a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.


```



```

a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.

```



```

a :- b, f.
b :- c, d.
b :- d, !, e.
b :- c.
c.
d.
e.

```

The query fails.

1.5.2 Applications of Cut

Cut can be used to

- improve efficiency
- force termination
- prevent duplicate solutions

But cuts can also change the meaning of a term, and so must be used with care.

Consider the factorial relation:

```

fac(0, 1).
fac(N, F) :- N1 is N-1,
              fac(N1, F1),
              F is F1*N.

```

This works OK for `fac(4, A)`, but hangs up on `fac(4, A)`, `A=10`

- `A=10` forces redo of `fac(0, X)`
- on redo, tries to use 2nd rule to take `0*R`, where `R` is found by `fac(-1, R)`.

Solution is to realize that any time we use the first rule, we don't want to redo with the second:

- The conditions for taking the two rules are mutually exclusive.

So add a cut:

```

fac(0, 1) :- !.
fac(N, F) :- N1 is N-1,
              fac(N1, F1),
              F is F1*N.

```

Consider a relation for checking to see if one list is a subset of another:

```

subset([], _).
subset([X|A], B) :- member(X, B),
                    subset(A, B).

```

What happens if this is invoked in a query:

```

L = [0,0,0,0,0, 1,1,1,1,1,
      2,2,2,2,2, 3,3,3,3,3,
      \ldots 9,9,9,9,9],
subset([1,2,3,4,5,6,7,8,9, a], L).

```

- We will get 9 successful member checks, then fail on member(a, L)
 - On redo, the next most recent member(9, L) succeeds again.
 - We fail again on member(a, L)
 - On redo, the next most recent member(9, L) succeeds again.
- ⋮
-

One way to speed this up is to use a cut:

```
subset2([], _).
subset2([X|A], B) :- member(X, B), !,
                    subset2(A, B).
```

This works very well when both operands are grounded, but not with a query like:

```
subset2([23], L).
```

What does the following predicate do?

```
foo(X) :- X, fail.
foo(X).
```

Now let's add a single cut:

What does the following predicate do?

```
foo(X) :- X, !, fail.
foo(X).
```

The combination `!, fail` is a common Prolog idiom for aborting a set of rules.

2 Programming in Prolog

1. Algorithmic Styles
2. Dare to Fail!
3. Open Lists

2.1 Algorithmic Styles

1. Backtracking
2. Guess and Verify

2.1.1 Backtracking

Prolog lends itself well to backtracking algorithms:

Given a problem with N variables X_1, \dots, X_N , each of which can take on one of k values: v_1, v_2, \dots, v_k , do

```
backtrack(i)
begin
  if (i > N) then
    print X's
    abort;
  else
    for j := 1 to k loop
      X[i] := v[j];
      backtrack(i+1);
    end loop
  end if
end;

backtrack(1);
```

In Prolog, this looks like

```
val(v1).
val(v2).
⋮
backtrack([]).
backtrack([X|L]) :-
    val(X), backtrack(L).
```

with the query:

```
backtrack([X1, X2, X3, \ldots]).
```

Problem is, backtracking is $O(k^N)$.

This makes Prolog one of the worlds easiest languages for writing exponential-time algorithms!

2.1.2 Guess and Verify

A better strategy in Prolog is to use what information you have to make a reasonable guess about the value of some variable, then check to see if that guess is correct.

A good example is

```
mustTakeBefore(C,D) :- prereq(C,D).
mustTakeBefore(C,D) :-
    prereq(C,E), mustTakeBefore(E,D).
```

- The `prereq` in the second rule is a “guess”, based upon the available list of facts.
- The remaining `mustTakeBefore` determines whether that guess is valid.

Note that this is much better than

```
mustTakeBefore(C,D) :- prereq(C,D).
mustTakeBefore(C,D) :-
    mustTakeBefore(E,D), prereq(C,E).
```

A characteristic of a good “guess” is that it should have fine (preferably small) possible solutions.

2.2 Dare to Fail!

We’ve already seen failure used in some creative ways to induce iteration.

Consider this as a way of generating guesses for a generate-and-test algorithm:

```
generate(0).
generate(X) :- generate(Y), X is Y+1.

squares :- generate(I),
            J is I*I,
            write(J), nl,
            I > 99.
```

This prints the first 101 integer squares. Why?

2.3 Open Lists

Sometimes it’s useful to have data structures that have variables in them as “holes” to be filled in later.

An example is the **open list**, a list ending in a variable.

```
L = [ a, b | X ].
L = [ a, b | X ], X = [ c | Y ].
```

- Note how the binding of `X` effectively appends to `L`, something that normally we can’t do.

- It’s almost like an assignment to an existing variable.
-

2.3.1 A Queue in Prolog

We can implement queues in Prolog as `q(X,Y)` where `X` is an open list and `Y` is the final variable in that open list.

```
setup(q(X,X)).

enter(A, q(X,Y), q(X,Z)) :- Y = [A|Z].

front(A, q(X,X)) :- !, fail.
front(A, q([A|_],Y)).

leave(q(X,X), _) :- !, fail.
leave(q(_|X],Y), q(X,Y)).
```

Sample use:

```
setup(Q), enter(a, Q, Q2),
enter(b, Q2, Q3), front(X, Q3).
```

3 Closing thoughts

- Prolog offers an interesting model that works in some very unconventional situations
- Does not scale up well to large projects.
- Very useful in contexts where some/all of the code will be machine-generated.
 - examples: software testing
- Prolog variants add pure numerics, reduce order dependencies, etc.